

Class 09 Structural Bioinformatics (Pt. 1)

Jiayi Zhou (PID:A17856751)

Table of contents

Exploring PDB Structures	4
PDB objects in R	8
Predict Protein Flexibility	16
Comparative structure analysis of Adenylate Kinase	19
Search and retrieve ADK structures	19

```
stats <- read.csv("Data Export Summary.csv")
stats
```

	Molecular.Type	X.ray	EM	NMR	Integrative	Multiple.methods
1	Protein (only)	176,378	20,438	12,709	342	221
2	Protein/Oligosaccharide	10,284	3,396	34	8	11
3	Protein/NA	9,007	5,931	287	24	7
4	Nucleic acid (only)	3,077	200	1,554	2	15
5	Other	174	13	33	3	0
6	Oligosaccharide (only)	11	0	6	0	1
	Neutron	Other	Total			
1	83	32	210,203			
2	1	0	13,734			
3	0	0	15,256			
4	3	1	4,852			
5	0	0	223			
6	0	4	22			

```
stats$Total
```

```
[1] "210,203" "13,734" "15,256" "4,852" "223" "22"
```

Oh, these are characters no numeric...

```
as.numeric(gsub(",","",stats$Total))
```

```
[1] 210203 13734 15256 4852    223     22
```

```
library(readr)
```

```
stats <- read_csv("Data Export Summary.csv")
```

```
Rows: 6 Columns: 9
-- Column specification -----
Delimiter: ","
chr (1): Molecular Type
dbl (4): Integrative, Multiple methods, Neutron, Other
num (4): X-ray, EM, NMR, Total

i Use `spec()` to retrieve the full column specification for this data.
i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
stats
```

```
# A tibble: 6 x 9
`Molecular Type` `X-ray`   EM   NMR Integrative `Multiple methods` Neutron
<chr>            <dbl> <dbl> <dbl>      <dbl>             <dbl> <dbl>
1 Protein (only) 176378 20438 12709      342           221    83
2 Protein/Oligosacch~ 10284  3396   34        8           11     1
3 Protein/NA       9007   5931   287       24            7     0
4 Nucleic acid (only) 3077   200   1554       2           15     3
5 Other            174    13    33        3            0     0
6 Oligosaccharide (o~ 11     0     6        0            1     0
# i 2 more variables: Other <dbl>, Total <dbl>
```

```
n.total <- sum(stats$Total)
n.total
```

```
[1] 244290
```

Q1: What percentage of structures in the PDB are solved by X-Ray and Electron Microscopy. Give your answer to 2 significant figures.

```
n.xray <- sum(stats$`X-ray`)
round(n.xray/ n.total*100)
```

[1] 81

```
n.xray <- sum(stats$`X-ray`)
percent.xray <- n.xray / n.total * 100
percent.xray
```

[1] 81.43231

81.43% of structures in the PDB are solved by X-Ray and Electron Microscopy.

There are 81.43 percent Xray structures in the PDB

Q2: What proportion of structures in the PDB are protein?

```
round(stats$Total[1]/n.total*100, 2)
```

[1] 86.05

```
library(readr)

stats <- read_csv("Data Export Summary.csv")
stats
```

```
# A tibble: 6 x 9
  `Molecular Type`    `X-ray`     EM     NMR Integrative `Multiple methods` Neutron
  <chr>              <dbl>    <dbl>   <dbl>      <dbl>          <dbl>    <dbl>
1 Protein (only)    176378  20438  12709      342        221     83
2 Protein/Oligosacch~ 10284   3396    34         8        11      1
3 Protein/NA        9007    5931   287       24         7      0
4 Nucleic acid (only) 3077    200   1554       2        15      3
5 Other             174     13    33         3         0      0
6 Oligosaccharide (o~ 11      0     6         0         1      0
# i 2 more variables: Other <dbl>, Total <dbl>
```

86.05 of structures in the PDB are protein.

Q3: Type HIV in the PDB website search box on the home page and determine how many HIV-1 protease structures are in the current PDB?

Current HIV-1 protease structures in the PDB: 873

Exploring PDB Structures

Package for structural bioinformatics

```
library(bio3d)  
  
hiv <- read.pdb("1hsg")
```

Note: Accessing on-line PDB file

```
hiv
```

Call: read.pdb(file = "1hsg")

Total Models#: 1
Total Atoms#: 1686, XYZs#: 5058 Chains#: 2 (values: A B)

Protein Atoms#: 1514 (residues/Calpha atoms#: 198)
Nucleic acid Atoms#: 0 (residues/phosphate atoms#: 0)

Non-protein/nucleic Atoms#: 172 (residues: 128)
Non-protein/nucleic resid values: [HOH (127), MK1 (1)]

Protein sequence:

```
PQITLWQRPLVTIKIGGQLKEALLDTGADDTVLEEMSLPGRWPKMIGGIGGFVKVRQYD  
QILIEICGHKAIGTVLVGPTPVNIIGRNLLTQIGCTLNFPQITLWQRPLVTIKIGGQLKE  
ALLDTGADDTVLEEMSLPGRWPKMIGGIGGFVKVRQYDQILIEICGHKAIGTVLVGPTP  
VNIIGRNLLTQIGCTLNF
```

```
+ attr: atom, xyz, seqres, helix, sheet,  
       calpha, remark, call
```

Let's first use the Mol* viewer to explore this structure.



Figure 1: My first view of HIV-Pr

```
hiv <- read.pdb("1hsg")
```

Note: Accessing on-line PDB file

Warning in get.pdb(file, path = tempdir(), verbose = FALSE):
/var/folders/hw(mb7lcvrd0v17c4nnmhd27vr40000gn/T//Rtmp1MpHUG/1hsg.pdb exists.
Skipping download

```
hiv
```

Call: read.pdb(file = "1hsg")

Total Models#: 1

Total Atoms#: 1686, XYZs#: 5058 Chains#: 2 (values: A B)

Protein Atoms#: 1514 (residues/Calpha atoms#: 198)

Nucleic acid Atoms#: 0 (residues/phosphate atoms#: 0)

Non-protein/nucleic Atoms#: 172 (residues: 128)

Non-protein/nucleic resid values: [HOH (127), MK1 (1)]

Protein sequence:

```
PQITLWQRPLVTIKIGGQLKEALLDTGADDTVLEEMSLPGRWPKMIGGIGGFIVKRQYD  
QILIEICGHKAIGTVLVGPTPVNIIGRNLLTQIGCTLNFPQITLWQRPLVTIKIGGQLKE  
ALLDTGADDTVLEEMSLPGRWPKMIGGIGGFIVKRQYDQILIEICGHKAIGTVLVGPTP  
VNIIGRNLLTQIGCTLNF
```

+ attr: atom, xyz, seqres, helix, sheet,
calpha, remark, call

Q4: Water molecules normally have 3 atoms. Why do we see just one atom per water molecule in this structure?


In X-ray crystal structures like HIV-1 protease (PDB: 1HSG), only the oxygen atom is visible and recorded in the PDB file because hydrogen atoms are not detected by X-ray crystallography, so each water appears as a single atom in the structure.

Q5: There is a critical “conserved” water molecule in the binding site. Can you identify this water molecule? What residue number does this water molecule have?

HOH 308 could be a conserved water molecule in the HIV-1 protease binding site because it is positioned closest to both the ligand and the catalytic Asp 25 residues, which is characteristic of the functionally important conserved water in this enzyme.

Q6: Generate and save a figure clearly showing the two distinct chains of HIV-protease along with the ligand. You might also consider showing the catalytic residues ASP 25 in each chain and the critical water (we recommend “Ball & Stick” for these side-chains). Add this figure to your Quarto document.

And a view of the ligand (ball and stick) with catalytic ASP 25 amino-acids (spacefill) and the all important active site water molecule (spacefill):



Indinavir and larger ligands or substrates can enter the HIV-1 protease binding site through the opening of flexible flaps formed by segments of each chain, which move apart to expose the catalytic site and then close to secure the ligand for catalysis or inhibition.

Q7. As you have hopefully observed HIV protease is a homodimer (i.e. it is composed of two identical chains). With the aid of the graphic display can you identify

secondary structure elements that are likely to only form in the dimer rather than the monomer?

In HIV protease, -sheets at the dimer interface are secondary structure elements that are likely to only form in the dimer and not in the monomer, as they are stabilized by interactions between the two identical chains visible in the graphic display

PDB objects in R

```
library(bio3d)
```

```
pdb <- read.pdb("1hsg")
```

Note: Accessing on-line PDB file

```
Warning in get.pdb(file, path = tempdir(), verbose = FALSE):
/var/folders/hw(mb7lcvrd0v17c4nnmhd27vr4000gn/T//Rtmp1MpHUG/1hsg.pdb exists.
Skipping download
```

```
pdb
```

Call: read.pdb(file = "1hsg")

Total Models#: 1

Total Atoms#: 1686, XYZs#: 5058 Chains#: 2 (values: A B)

Protein Atoms#: 1514 (residues/Calpha atoms#: 198)

Nucleic acid Atoms#: 0 (residues/phosphate atoms#: 0)

Non-protein/nucleic Atoms#: 172 (residues: 128)

Non-protein/nucleic resid values: [HOH (127), MK1 (1)]

Protein sequence:

```
PQITLWQRPLVTIKIGGQLKEALLDTGADDTVLEEMSLPGRWPKMIGGIGGFIKVRQYD  
QILIEICGHKAIGTVLVGPTPVNIIGRNLLTQIGCTLNFPQITLWQRPLVTIKIGGQLKE  
ALLDTGADDTVLEEMSLPGRWPKMIGGIGGFIKVRQYDQILIEICGHKAIGTVLVGPTP  
VNIIGRNLLTQIGCTLNF
```

```
+ attr: atom, xyz, seqres, helix, sheet,  
       calpha, remark, call
```

Q7: How many amino acid residues are there in this pdb object?

198 residues.

```
sum(pdb$calpha)
```

[1] 198

Q8: Name one of the two non-protein residues?

HOH (water) or MK1 (indinavir).

```
unique(pdb$atom$resid[pdb$atom$type == "HETATM"])
```

[1] "MK1" "HOH"

Q9: How many protein chains are in this structure?

2 chains (A and B).

```
unique(pdb$atom$chain[pdb$atom$type == "ATOM"])
```

[1] "A" "B"

```
length(unique(pdb$atom$chain[pdb$atom$type == "ATOM"]))
```

[1] 2

```
head(hiv$atom)
```

	type	eleno	elety	alt	resid	chain	resno	insert	x	y	z	o	b
1	ATOM	1	N	<NA>	PRO	A	1	<NA>	29.361	39.686	5.862	1	38.10
2	ATOM	2	CA	<NA>	PRO	A	1	<NA>	30.307	38.663	5.319	1	40.62
3	ATOM	3	C	<NA>	PRO	A	1	<NA>	29.760	38.071	4.022	1	42.64
4	ATOM	4	O	<NA>	PRO	A	1	<NA>	28.600	38.302	3.676	1	43.40
5	ATOM	5	CB	<NA>	PRO	A	1	<NA>	30.508	37.541	6.342	1	37.87
6	ATOM	6	CG	<NA>	PRO	A	1	<NA>	29.296	37.591	7.162	1	38.40
	segid	elesy	charge										
1	<NA>	N	<NA>										
2	<NA>	C	<NA>										
3	<NA>	C	<NA>										
4	<NA>	O	<NA>										
5	<NA>	C	<NA>										
6	<NA>	C	<NA>										

Extract the sequence

```
 pdbseq(hiv)
```

```
 1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20  
"P" "Q" "I" "T" "L" "W" "Q" "R" "P" "L" "V" "T" "I" "K" "I" "G" "G" "Q" "L" "K"  
21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40  
"E" "A" "L" "L" "D" "T" "G" "A" "D" "D" "T" "V" "L" "E" "E" "M" "S" "L" "P" "G"  
41 42 43 44 45 46 47 48 49 50 51 52 53 54 55 56 57 58 59 60  
"R" "W" "K" "P" "K" "M" "I" "G" "G" "I" "G" "F" "I" "K" "V" "R" "Q" "Y" "D"  
61 62 63 64 65 66 67 68 69 70 71 72 73 74 75 76 77 78 79 80  
"Q" "I" "L" "I" "E" "I" "C" "G" "H" "K" "A" "I" "G" "T" "V" "L" "V" "G" "P" "T"  
81 82 83 84 85 86 87 88 89 90 91 92 93 94 95 96 97 98 99 1  
"P" "V" "N" "I" "I" "G" "R" "N" "L" "L" "T" "Q" "I" "G" "C" "T" "L" "N" "F" "P"  
2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20 21  
"Q" "I" "T" "L" "W" "Q" "R" "P" "L" "V" "T" "I" "K" "I" "G" "G" "Q" "L" "K" "E"  
22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40 41  
"A" "L" "L" "D" "T" "G" "A" "D" "D" "T" "V" "L" "E" "E" "M" "S" "L" "P" "G" "R"  
42 43 44 45 46 47 48 49 50 51 52 53 54 55 56 57 58 59 60 61  
"W" "K" "P" "K" "M" "I" "G" "G" "I" "G" "F" "I" "K" "V" "R" "Q" "Y" "D" "Q"  
62 63 64 65 66 67 68 69 70 71 72 73 74 75 76 77 78 79 80 81  
"I" "L" "I" "E" "I" "C" "G" "H" "K" "A" "I" "G" "T" "V" "L" "V" "G" "P" "T" "P"  
82 83 84 85 86 87 88 89 90 91 92 93 94 95 96 97 98 99  
"V" "N" "I" "I" "G" "R" "N" "L" "L" "T" "Q" "I" "G" "C" "T" "L" "N" "F"
```

```
chainA_seq <- pdbseq(trim.pdb(hiv, chain="A"))
```

```
attributes(pdb)
```

```
$names  
[1] "atom"    "xyz"      "seqres"   "helix"    "sheet"    "calpha"   "remark"   "call"
```

```
$class  
[1] "pdb"    "sse"
```

```
head(pdb$atom)
```

	type	eleno	elety	alt	resid	chain	resno	insert	x	y	z	o	b
1	ATOM	1	N	<NA>	PRO	A	1	<NA>	29.361	39.686	5.862	1	38.10
2	ATOM	2	CA	<NA>	PRO	A	1	<NA>	30.307	38.663	5.319	1	40.62

```

3 ATOM    3      C <NA>   PRO     A     1    <NA> 29.760 38.071 4.022 1 42.64
4 ATOM    4      O <NA>   PRO     A     1    <NA> 28.600 38.302 3.676 1 43.40
5 ATOM    5      CB <NA>  PRO     A     1    <NA> 30.508 37.541 6.342 1 37.87
6 ATOM    6      CG <NA>  PRO     A     1    <NA> 29.296 37.591 7.162 1 38.40
segid elesy charge
1 <NA>      N  <NA>
2 <NA>      C  <NA>
3 <NA>      C  <NA>
4 <NA>      O  <NA>
5 <NA>      C  <NA>
6 <NA>      C  <NA>

```

I can interactively view these PDB objects in R with the new **bio3dview** package. This is not yet on CRAN.

To install this I can setup install **bio3dview pak** package and use it to run from Github. In my console I first run

```
#install.packages("pak")
pak::pak("bioboot/bio3dview")
```

```
! Using bundled GitHub PAT. Please add your own PAT using `gitcreds::gitcreds_set()`.
```

```
Loading metadata database
```

```
Loading metadata database ... done
```

```
No downloads are needed
```

```
1 pkg + 40 deps: kept 40 [7s]
```

```
#install.packages("NGLVieweR")
```

```
library(bio3dview)
```

```
view.pdb(hiv)
```

PhantomJS not found. You can install it with `webshot::install_phantomjs()`. If it is installed
file:///private/var/folders/hw/mb7lcvrd0v17c4nnmhd27vr40000gn/T/Rtmp1MpHUG/file2de02886d17a




Change some settings

```
# Select the important ASP 25 residue
sele <- atom.select(pdb, resno=25)

# and highlight them in spacefill representation
view.pdb(hiv, highlight = sele,
          highlight.style = "spacefill",
          colorScheme = "chain",
          col=c("navy", "orange"),
          backgroundColor = "lightyellow")
```

file:///private/var/folders/hw(mb71cvrd0v17c4nnmhd27vr40000gn/T/Rtmp1MpHUG/file2de059dc2e8a



Predict Protein Flexibility

We can run a bioinformatic calculation to predict protein dynamics - i.e. functional motions.

We will use the `nma()` function:

```
adk <- read.pdb("6s36")
```

Note: Accessing on-line PDB file
PDB has ALT records, taking A only, rm.alt=TRUE

```
adk
```

Call: `read.pdb(file = "6s36")`

Total Models#: 1

Total Atoms#: 1898, XYZs#: 5694 Chains#: 1 (values: A)

Protein Atoms#: 1654 (residues/Calpha atoms#: 214)

Nucleic acid Atoms#: 0 (residues/phosphate atoms#: 0)

Non-protein/nucleic Atoms#: 244 (residues: 244)

Non-protein/nucleic resid values: [CL (3), HOH (238), MG (2), NA (1)]

Protein sequence:

```
MRIILLGAPGAGKGTQAQFIMEKYGIPQISTGMLRAAVKSGSELGKQAKDIMDAGKLVT  
DELVIALVKERIAQEDCRNGFLLDGFPRТИQADAMKEAGINVDTVLEFDVPDELVDKI  
VGRRVHAPSGRVYHVKFNPPKVEGKDDVTGEELTRKDDQEETVRKRLVEYHQMTAPLIG  
YYSKAEAGNTKYAKVDGTPVAEVRADLEKILG
```


+ attr: atom, xyz, seqres, helix, sheet,
calpha, remark, call

```
m <- nma(adk)
```

Building Hessian... Done in 0.022 seconds.

Diagonalizing Hessian... Done in 0.417 seconds.

```
plot(m)
```




Generate a “trrajectory” of predicted motion

```
mktrj(m, file="ADK_nma.pdb")
```

```
view.nma(m)
```

```
file:///private/var/folders/hw/mb7lcvrd0v17c4nnmhd27vr40000gn/T/Rtmp1MpHUG/file2de03e888b79
```



Comparative structure analysis of Adenylate Kinase

```
# Install packages in the R console NOT your Rmd/Quarto file

#install.packages("bio3d")
#install.packages("NGLVieweR")

#install.packages("pak")
#pak::pak("bioboot/bio3dview")

#install.packages("BiocManager")
#BiocManager::install("msa")
```

Q10. Which of the packages above is found only on BioConductor and not CRAN?

msa. It's a Bioconductor package (installed via `BiocManager::install("msa")`), not on CRAN.

Q11. Which of the above packages is not found on BioConductor or CRAN?:

bio3dview. It's GitHub-only (`pak::pak("bioboot/bio3dview")`), not on CRAN or Bioconductor.

Q12. True or False? Functions from the pak package can be used to install packages from GitHub and BitBucket?

True. pak can install from GitHub and Bitbucket using remote specs (e.g., “user/repo” for GitHub, `bitbucket::user/repo`).

Search and retrieve ADK structures

```
library(bio3d)
aa <- get.seq("1ake_A")
```

Warning in get.seq("1ake_A"): Removing existing file: seqs.fasta

Fetching... Please wait. Done.

aa

	1	60
pdb 1AKE A	MRIILLGAPGAGKGTQAQFIMEKYGIPQISTGMLRAAVKSGSELGKQAKDIDAGKLVT					
	1	60
	61	120
pdb 1AKE A	DELVIALVKERIAQEDCRNGFLLDGFPRTIPQADAMKEAGINVDYVLEFDVPDELIVDRI					
	61	120
	121	180
pdb 1AKE A	VGRRVHAPSGRKYHVKFNPPKVEGKDDVTGEELTRKDDQEETVRKRLVEYHQMTAPLIG					
	121	180
	181	.	.	.	214	
pdb 1AKE A	YYSKAEAGNTKYAKVADGTPVAEVRADLEKILG					
	181	.	.	.	214	

Call:

```
read.fasta(file = outfile)
```

Class:

```
fasta
```

Alignment dimensions:

```
1 sequence rows; 214 position columns (214 non-gap, 0 gap)
```

```
+ attr: id, ali, call
```

Q13. How many amino acids are in this sequence, i.e. how long is this sequence?

214 amino acids.