

Class 10: Halloween Mini Project

Jiayi Zhou (PID:A17856751)

Table of contents

Data Import	1
Quick Overview of the dataset	4
Winpercent and Pricepercent	14
Exploring the correlation structure	17
Principle Componenet Analysis	19

As it is nearly Halloween and the half way point in the quarter let's do a mini project to help us figure out the best candy!

Our come from the 538 website and is available as a CSV file:

Data Import

```
candy <- read.csv("candy-data.csv", row.names = 1)
head(candy)
```

	chocolate	fruity	caramel	peanut	yalmond	nougat	crisped	rice	wafer
100 Grand	1	0	1		0	0			1
3 Musketeers	1	0	0		0	1			0
One dime	0	0	0		0	0			0
One quarter	0	0	0		0	0			0
Air Heads	0	1	0		0	0			0
Almond Joy	1	0	0		1	0			0
	hard	bar	pluribus	sugarpercent	pricepercent	winpercent			
100 Grand	0	1	0	0.732	0.860	66.97173			
3 Musketeers	0	1	0	0.604	0.511	67.60294			
One dime	0	0	0	0.011	0.116	32.26109			

One quarter	0	0	0	0.011	0.511	46.11650
Air Heads	0	0	0	0.906	0.511	52.34146
Almond Joy	0	1	0	0.465	0.767	50.34755

```
flextable::flextable(head(candy, 10))
```

chocolate	fruity	caramel	peanut	almond	nougat	crispedrice	wafer	hard	bar	pluribus	s
1	0	1	0	0	1	0	0	1	1	0	
1	0	0	0	0	1	0	0	0	1	0	
0	0	0	0	0	0	0	0	0	0	0	
0	0	0	0	0	0	0	0	0	0	0	
0	1	0	0	0	0	0	0	0	0	0	
1	0	0	1	0	0	0	0	0	1	0	
1	0	1	1	1	0	0	0	0	1	0	
0	0	0	1	0	0	0	0	0	0	1	
0	0	0	0	0	0	0	0	0	0	1	
0	1	1	0	0	0	0	0	0	0	0	

Q1. How many different candy types are in this dataset?

```
nrow(candy)
```

```
[1] 85
```

```
library(tidyverse)
```

```
-- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
v dplyr     1.1.4     v readr     2.1.5
v forcats   1.0.1     v stringr   1.5.2
v ggplot2   4.0.0     v tibble    3.3.0
v lubridate  1.9.4     v tidyr    1.3.1
v purrr     1.1.0
-- Conflicts ----- tidyverse_conflicts() --
x dplyr::filter() masks stats::filter()
x dplyr::lag()   masks stats::lag()
i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to becom
```

```
candy %>%
  nrow()
```

[1] 85

Q2. How many fruity candy types are in the dataset?

```
sum(candy$fruity)
```

[1] 38

Q3. What is your favorite candy in the dataset and what is it's winpercent value?

My favorite winpercent

```
candy["Hershey's Milk Chocolate", ]$winpercent
```

[1] 56.4905

```
library(dplyr)

candy |>
  filter(rownames(candy) == "Hershey's Milk Chocolate") |>
  select(winpercent)
```

	winpercent
Hershey's Milk Chocolate	56.4905

Q4. What is the winpercent value for “Kit Kat”?

```
candy |>
  filter(rownames(candy) == "Kit Kat") |>
  select(winpercent)
```

	winpercent
Kit Kat	76.7686

Q5. What is the winpercent value for “Tootsie Roll Snack Bars”?

```
candy |>
  filter(rownames(candy) == "Tootsie Roll Snack Bars") |>
  select(winpercent)
```

```
winpercent
Tootsie Roll Snack Bars    49.6535
```

Quick Overview of the dataset

```
library("skimr")
skimr::skim(candy)
```

Table 2: Data summary

Name	candy
Number of rows	85
Number of columns	12
Column type frequency:	
numeric	12
Group variables	
	None

Variable type: numeric

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
chocolate	0	1	0.44	0.50	0.00	0.00	0.00	1.00	1.00	
fruity	0	1	0.45	0.50	0.00	0.00	0.00	1.00	1.00	
caramel	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
peanutyalmond	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
nougat	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
crispedricewafer	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
hard	0	1	0.18	0.38	0.00	0.00	0.00	0.00	1.00	
bar	0	1	0.25	0.43	0.00	0.00	0.00	0.00	1.00	
pluribus	0	1	0.52	0.50	0.00	0.00	1.00	1.00	1.00	
sugarpercent	0	1	0.48	0.28	0.01	0.22	0.47	0.73	0.99	
pricepercent	0	1	0.47	0.29	0.01	0.26	0.47	0.65	0.98	
winpercent	0	1	50.32	14.71	22.45	39.14	47.83	59.86	84.18	

Q6. Is there any variable/column that looks to be on a different scale to the majority of the other columns in the dataset?

The winpercent is on 0-100 scale and the rest are 0-1 scale.

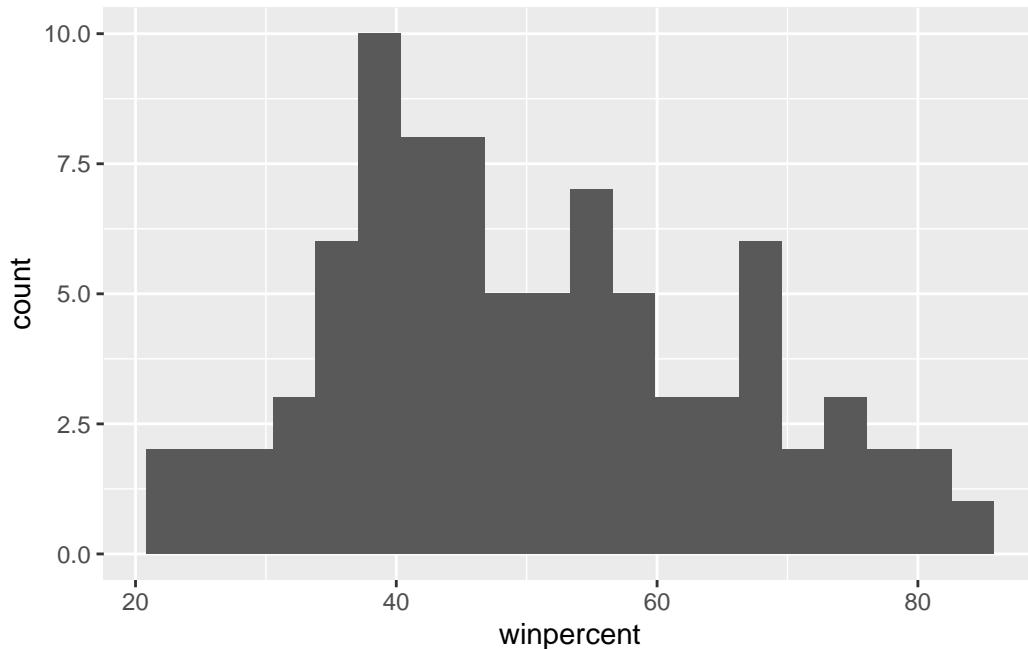
Q7. What do you think a zero and one represent for the candy\$chocolate column?

That the candy does not contain chocolate.

Q8. Plot a histogram of winpercent values

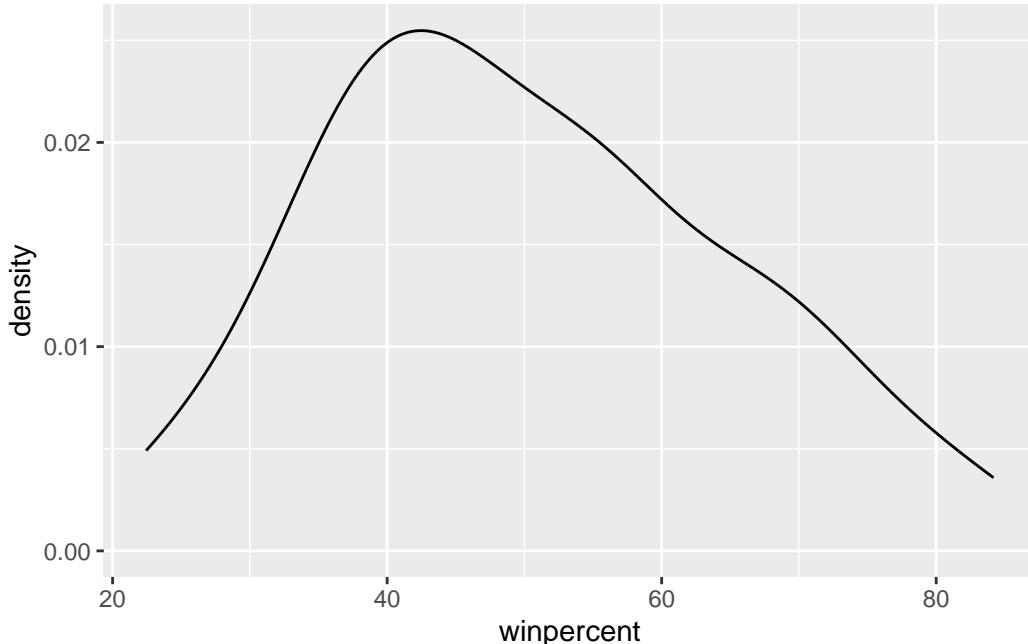
```
library(ggplot2)

ggplot(candy) +
  aes(winpercent) +
  geom_histogram(bins=20)
```



```
library(ggplot2)

ggplot(candy) +
  aes(winpercent) +
  geom_density()
```



Q9. Is the distribution of winpercent values symmetrical?

Based on the summary statistics, the mean (50.32) is greater than the median (47.83), and the maximum (84.18) is much farther from the median than the minimum (22.45). The distribution is not symmetrical; it is right-skewed (positively skewed)

Q10. Is the center of the distribution above or below 50%?

```
mean(candy$winpercent)
```

[1] 50.31676

```
summary(candy$winpercent)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
22.45	39.14	47.83	50.32	59.86	84.18

The median winpercent is 47.83, which is below 50%. Therefore, the center of the distribution is below 50%

Q11. On average is chocolate candy higher or lower ranked than fruit candy?

Chocolate is higher ranked on average.

```
# 1. Find all chocolate candy in the dataset.  
# 2, Find their winpercent values  
# Calculate the mean of these values  
  
# 4-6. Do the same for fruity candy  
# 7. Compare mean winpercents of chocolate vs fruity  
# 8. Pick the highest as the winner  
  
choc.inds <- candy$chocolate==1  
choc.win <- candy[choc.inds,]$winpercent  
choc.win <- mean(choc.win)  
choc.win
```

```
[1] 60.92153
```

```
mean(candy[candy$chocolate==1,]$winpercent)
```

```
[1] 60.92153
```

```
mean(candy[candy$fruity==1,]$winpercent)
```

```
[1] 44.11974
```

```
fruit.ind <- candy$fruity==1  
fruit.win <- candy[fruit.ind,]$winpercent  
fruit.mean <- mean(fruit.win)  
  
choc.ind <- candy$chocolate==1  
choc.win <- candy[choc.ind,]$winpercent  
choc.mean <- mean(choc.win)
```

```
candy |>  
  filter(chocolate == 1) |>  
  select(winpercent)
```

	winpercent
100 Grand	66.97173

3 Musketeers	67.60294
Almond Joy	50.34755
Baby Ruth	56.91455
Charleston Chew	38.97504
Hershey's Kisses	55.37545
Hershey's Krackel	62.28448
Hershey's Milk Chocolate	56.49050
Hershey's Special Dark	59.23612
Junior Mints	57.21925
Kit Kat	76.76860
Peanut butter M&M's	71.46505
M&M's	66.57458
Milk Duds	55.06407
Milky Way	73.09956
Milky Way Midnight	60.80070
Milky Way Simply Caramel	64.35334
Mounds	47.82975
Mr Good Bar	54.52645
Nestle Butterfinger	70.73564
Nestle Crunch	66.47068
Peanut M&Ms	69.48379
Reese's Miniatures	81.86626
Reese's Peanut Butter cup	84.18029
Reese's pieces	73.43499
Reese's stuffed with pieces	72.88790
Rolo	65.71629
Sixlets	34.72200
Nestle Smarties	37.88719
Snickers	76.67378
Snickers Crisper	59.52925
Tootsie Pop	48.98265
Tootsie Roll Juniors	43.06890
Tootsie Roll Midgies	45.73675
Tootsie Roll Snack Bars	49.65350
Twix	81.64291
Whoppers	49.52411

Q12. Is this difference statistically significant?

Yes, the difference is statistically significant.

```
t.test(choc.win, fruit.win)
```

Welch Two Sample t-test

```
data: choc.win and fruit.win
t = 6.2582, df = 68.882, p-value = 2.871e-08
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 11.44563 22.15795
sample estimates:
mean of x mean of y
60.92153 44.11974
```

Q13. What are the five least liked candy types in this set?

```
candy |>
  arrange(winpercent) |>
  head(5)
```

	chocolate	fruity	caramel	peanuty	almondy	nougat
Nik L Nip	0	1	0	0	0	0
Boston Baked Beans	0	0	0	1	0	0
Chiclets	0	1	0	0	0	0
Super Bubble	0	1	0	0	0	0
Jawbusters	0	1	0	0	0	0
	crispedrice	wafers	hard bar	pluribus	sugarpercent	pricepercent
Nik L Nip	0	0	0	1	0.197	0.976
Boston Baked Beans	0	0	0	1	0.313	0.511
Chiclets	0	0	0	1	0.046	0.325
Super Bubble	0	0	0	0	0.162	0.116
Jawbusters	0	1	0	1	0.093	0.511
	winpercent					
Nik L Nip	22.44534					
Boston Baked Beans	23.41782					
Chiclets	24.52499					
Super Bubble	27.30386					
Jawbusters	28.12744					

```
x <- c(5,1,10,4)
sort(x)
```

[1] 1 4 5 10

```
#sort(x)
#(candy$wonpercent)
```

```
ord.ind <- order(candy$winpercent)
head(candy[ord.ind,],5)
```

	chocolate	fruity	caramel	peanuty	almondy	nougat	
Nik L Nip	0	1	0		0	0	
Boston Baked Beans	0	0	0		1	0	
Chiclets	0	1	0		0	0	
Super Bubble	0	1	0		0	0	
Jawbusters	0	1	0		0	0	
	crispedrice	wafers	hard	bar	pluribus	sugarpercent	pricepercent
Nik L Nip	0	0	0	1		0.197	0.976
Boston Baked Beans	0	0	0	1		0.313	0.511
Chiclets	0	0	0	1		0.046	0.325
Super Bubble	0	0	0	0		0.162	0.116
Jawbusters	0	1	0	1		0.093	0.511
	winpercent						
Nik L Nip	22.44534						
Boston Baked Beans	23.41782						
Chiclets	24.52499						
Super Bubble	27.30386						
Jawbusters	28.12744						

Q14. What are the top 5 all time favorite candy types out of this set?

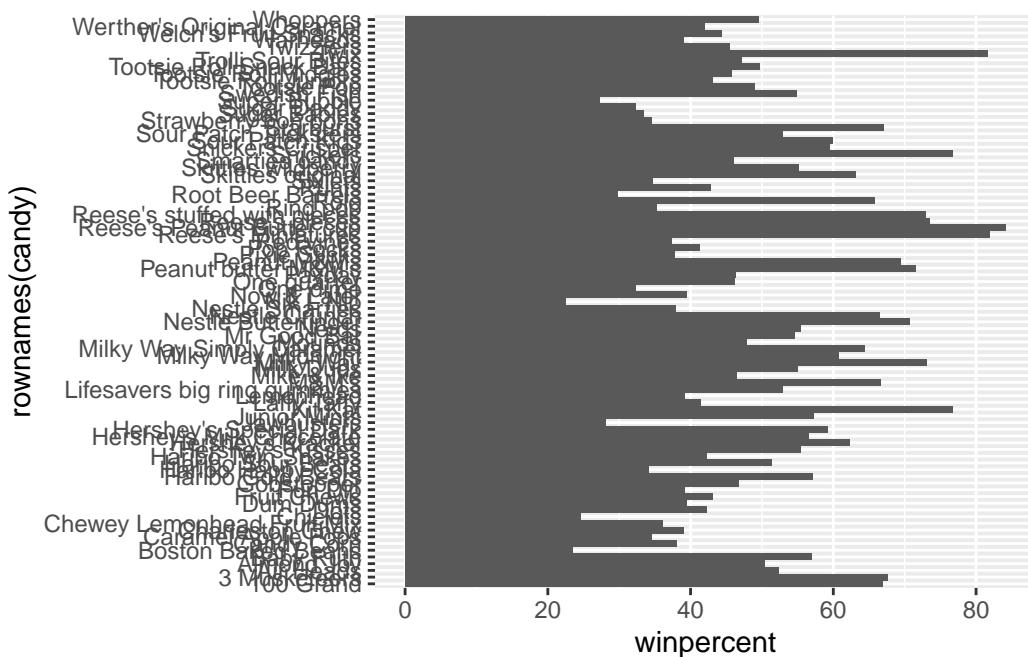
```
candy |>
  arrange(-winpercent) |>
  head(5)
```

	chocolate	fruity	caramel	peanuty	almondy	nougat
Reese's Peanut Butter cup	1	0	0		1	0
Reese's Miniatures	1	0	0		1	0
Twix	1	0	1		0	0
Kit Kat	1	0	0		0	0
Snickers	1	0	1		1	1
	crispedrice	wafers	hard	bar	pluribus	sugarpercent
Reese's Peanut Butter cup	0	0	0	0		0.720
Reese's Miniatures	0	0	0	0		0.034
Twix	1	0	1	0		0.546

	price	percent	win	percent	
Kit Kat	1	0	1	0	0.313
Snickers	0	0	1	0	0.546
Reese's Peanut Butter cup	0.651	84.18029			
Reese's Miniatures	0.279	81.86626			
Twix	0.906	81.64291			
Kit Kat	0.511	76.76860			
Snickers	0.651	76.67378			

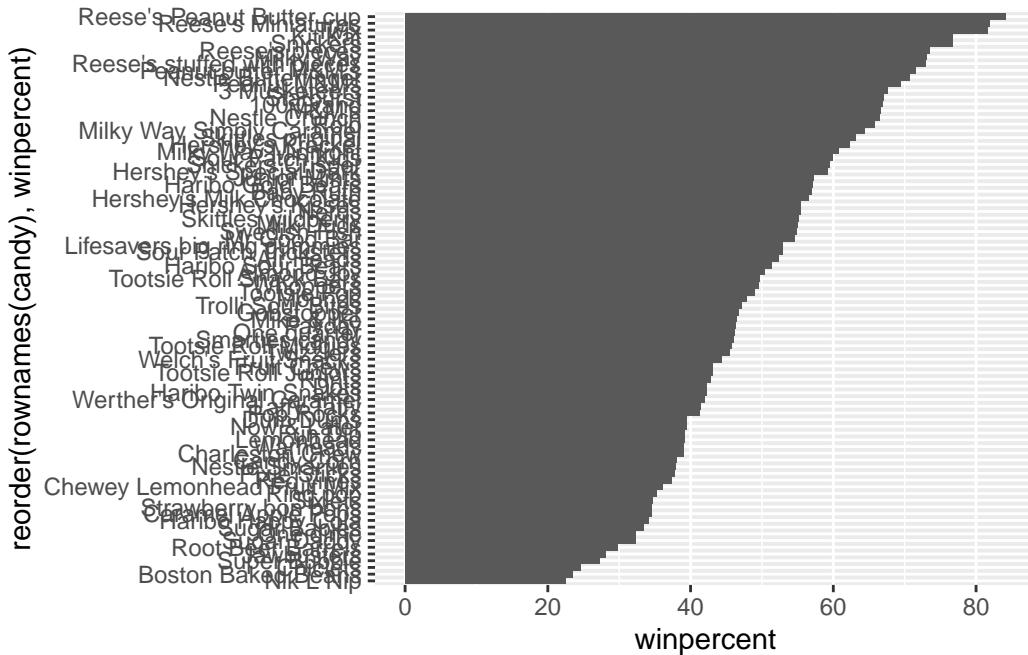
Q15. Make a first barplot of candy ranking based on winpercent values.

```
ggplot(candy) +  
  aes(winpercent, rownames(candy)) +  
  geom_col()
```



Q16. This is quite ugly, use the reorder() function to get the bars sorted by winpercent?

```
ggplot(candy) +  
  aes(x=winpercent,  
       y=reorder(rownames(candy), winpercent)) +  
  geom_col()
```

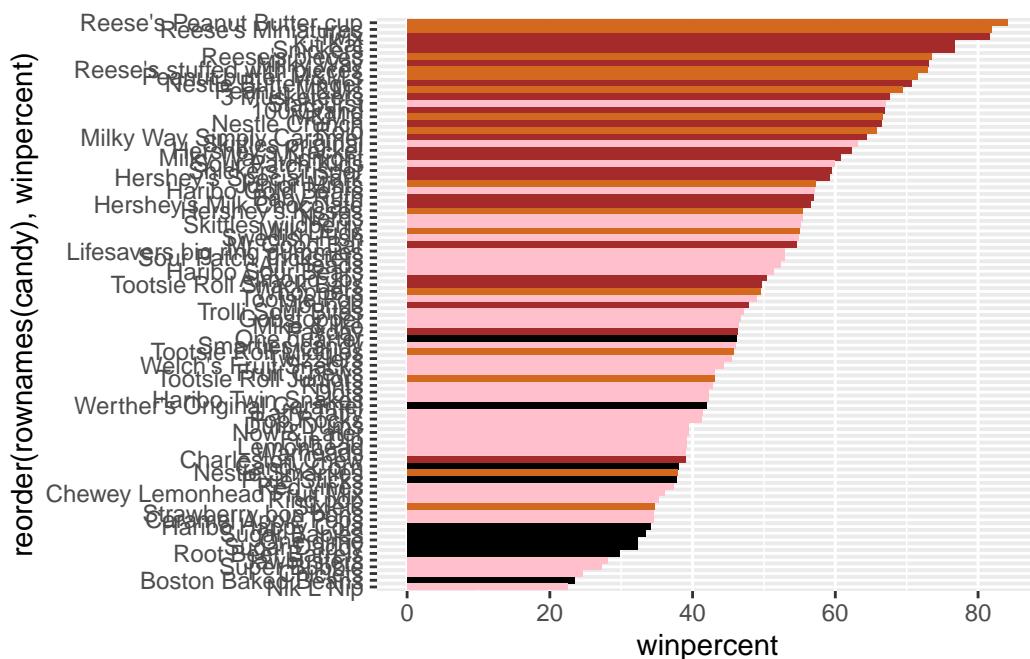


Add some colors based on the “type of candy”

```
my_cols=rep("black", nrow(candy))
my_cols[as.logical(candy$chocolate)] = "chocolate"
my_cols[as.logical(candy$fruity)] = "pink"
my_cols[as.logical(candy$bar)] = "brown"
my_cols
```

```
[1] "brown"      "brown"       "black"        "black"        "pink"         "brown"
[7] "brown"      "black"       "black"        "pink"         "brown"        "pink"
[13] "pink"        "pink"        "pink"         "pink"         "pink"         "pink"
[19] "pink"        "black"       "pink"         "pink"         "chocolate"    "brown"
[25] "brown"       "brown"       "pink"         "chocolate"    "brown"        "pink"
[31] "pink"        "pink"        "chocolate"    "chocolate"    "pink"         "chocolate"
[37] "brown"       "brown"       "brown"        "brown"        "brown"        "pink"
[43] "brown"       "brown"       "pink"         "pink"         "brown"        "chocolate"
[49] "black"        "pink"        "pink"         "chocolate"    "chocolate"    "chocolate"
[55] "chocolate"   "pink"        "chocolate"    "black"        "pink"         "chocolate"
[61] "pink"        "pink"        "chocolate"    "pink"         "brown"        "brown"
[67] "pink"        "pink"        "pink"         "pink"         "black"        "black"
[73] "pink"        "pink"        "pink"         "chocolate"    "chocolate"    "brown"
[79] "pink"        "brown"       "pink"         "pink"         "pink"         "black"
[85] "chocolate"
```

```
ggplot(candy) +
  aes(x=winpercent,
      y=reorder(rownames(candy), winpercent)) +
  geom_col(fill=my_cols)
```



Now, for the first time, using this plot we can answer questions like:

Q17. What is the worst ranked chocolate candy?

```
candy |>
  dplyr::filter(chocolate == 1) |>
  dplyr::slice_min(winpercent, n = 1, with_ties = FALSE)
```

	chocolate	fruity	caramel	peanut	almond	nougat	crispedrice	wafer	hard
Sixlets	1	0	0	0	0	0	0	0	0
	bar	pluribus	sugar	percent	price	percent	win	percent	
Sixlets	0	1	0.22	0.081	34.722				

Sixlets.

Q18. What is the best ranked fruity candy?

```
candy |>  
  dplyr::filter(fruity == 1) |>  
  dplyr::slice_max(winpercent, n = 1, with_ties = FALSE)
```

	chocolate	fruity	caramel	peanuty	almondy	nougat	crispedrice	wafer	hard
Starburst	0	1	0			0	0		0
	bar	pluribus	sugar	percent	price	percent	win	percent	
Starburst	0	1	0.151		0.22	67.03763			

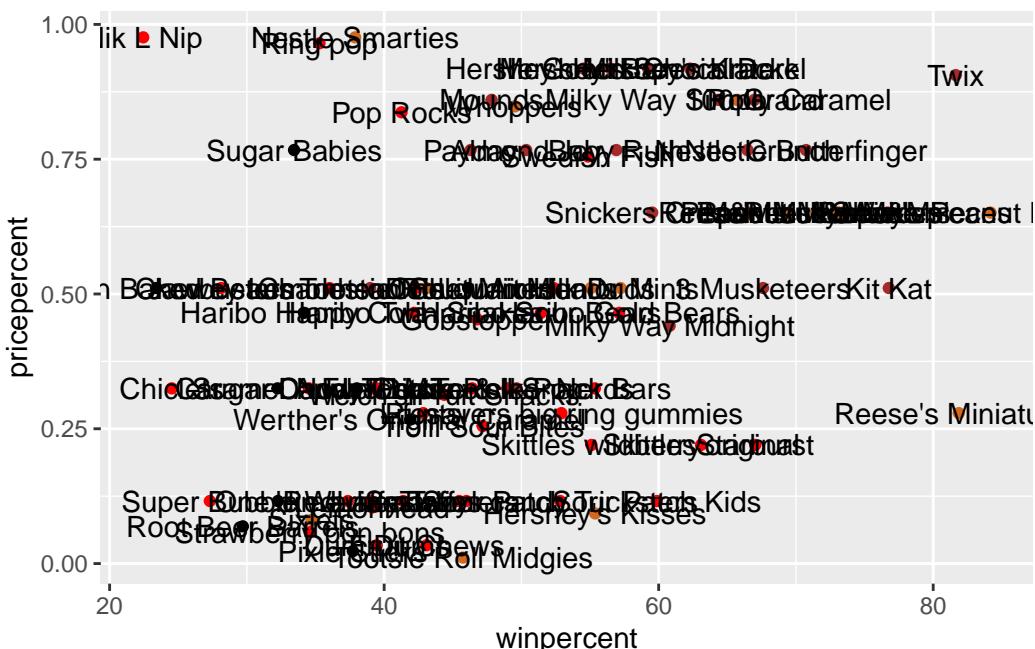
Starburst.

Winpercent and Pricepercent

A plot with both variables/columns winpercent and pricepercent

```
my_cols[as.logical(candy$fruity)] = "red"

ggplot(candy) +
  aes(x=winpercent, y=pricepercent, label=row.names(candy)) +
  geom_point(col=my_cols) +
  geom_text()
```



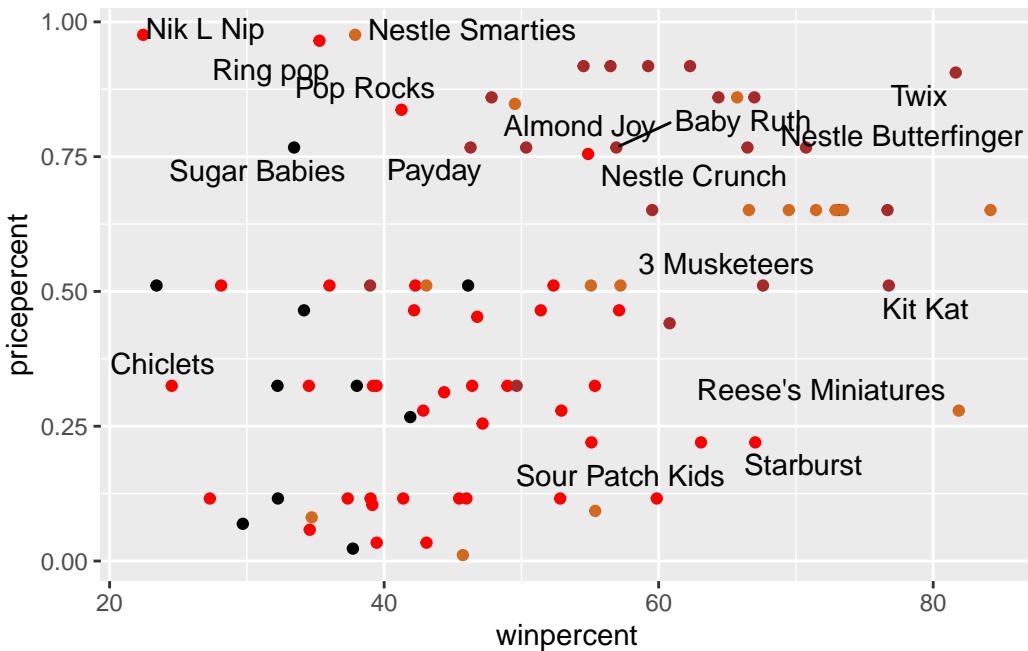
```

library(ggrepel)

ggplot(candy) +
  aes(x=winpercent, y=pricepercent, label=row.names(candy)) +
  geom_point(col=my_cols) +
  geom_text_repel(max.overlaps = 7)

```

Warning: ggrepel: 68 unlabeled data points (too many overlaps). Consider increasing max.overlaps



Q19. Which candy type is the highest ranked in terms of winpercent for the least money - i.e. offers the most bang for your buck?

```

candy |>
  dplyr::mutate(type = dplyr::case_when(chocolate == 1 & bar == 1 ~ "chocolate bar",
                                            chocolate == 1 & bar == 0 ~ "chocolate other",
                                            fruity == 1 ~ "fruity",
                                            TRUE ~ "other")) |>
  dplyr::group_by(type) |>
  dplyr::summarise(mean_win = mean(winpercent, na.rm = TRUE),
                  mean_price = mean(pricepercent, na.rm = TRUE), .groups = "drop") |>

```

```
dplyr::mutate(score = dplyr::min_rank(mean_price) + dplyr::min_rank(dplyr::desc(mean_win)))
dplyr::slice_min(score, n = 1, with_ties = FALSE)
```

```
# A tibble: 1 x 4
  type    mean_win mean_price score
  <chr>     <dbl>      <dbl> <int>
1 fruity      44.0       0.333     4
```

Fruity type is the highest ranked in terms of winpercent for the least money.

```
candy |>
  dplyr::mutate(score = dplyr::min_rank(pricepercent) +
                 dplyr::min_rank(dplyr::desc(winpercent))) |>
  dplyr::slice_min(score, n = 1, with_ties = FALSE)
```

	chocolate	fruity	caramel	peanuty	almondy	nougat
Reese's Miniatures	1	0	0		1	0
	crisp	pedri	cewafer	hard	bar	pluribus
Reese's Miniatures						sugarpercent
	winpercent	score				pricepercent
Reese's Miniatures	81.86626	26				0.279

Reese's Miniatures is the highest single candy ranked in terms of winpercent for the least money

Q20. What are the top 5 most expensive candy types in the dataset and of these which is the least popular?

```
candy |>
  arrange(-pricepercent) |>
  head(5)
```

	chocolate	fruity	caramel	peanuty	almondy	nougat
Nik L Nip	0	1	0		0	0
Nestle Smarties	1	0	0		0	0
Ring pop	0	1	0		0	0
Hershey's Krackel	1	0	0		0	0
Hershey's Milk Chocolate	1	0	0		0	0
	crisp	pedri	cewafer	hard	bar	pluribus
Nik L Nip	0	0	0		1	sugarpercent
						0.197

Nestle Smarties	0	0	0	1	0.267
Ring pop	0	1	0	0	0.732
Hershey's Krackel	1	0	1	0	0.430
Hershey's Milk Chocolate	0	0	1	0	0.430
	pricepercent		winpercent		
Nik L Nip	0.976	22.44534			
Nestle Smarties	0.976	37.88719			
Ring pop	0.965	35.29076			
Hershey's Krackel	0.918	62.28448			
Hershey's Milk Chocolate	0.918	56.49050			

```
candy |>
  arrange(-pricepercent) |>
  head(5) |>
  dplyr::slice_min(winpercent, n = 1, with_ties = FALSE)
```

	chocolate	fruity	caramel	peanuty	almondy	nougat	crispedrice	wafer	hard	
Nik L Nip	0	1	0			0	0		0	0
	bar	pluribus	sugarpercent	pricepercent	winpercent					
Nik L Nip	0	1	0.197	0.976	22.44534					

Nik L Nip, Ring Pop, Nestlé Smarties, Hershey's Krackel, Hershey's Milk Chocolate. Among these, the least popular (lowest winpercent) is Nik L Nip.

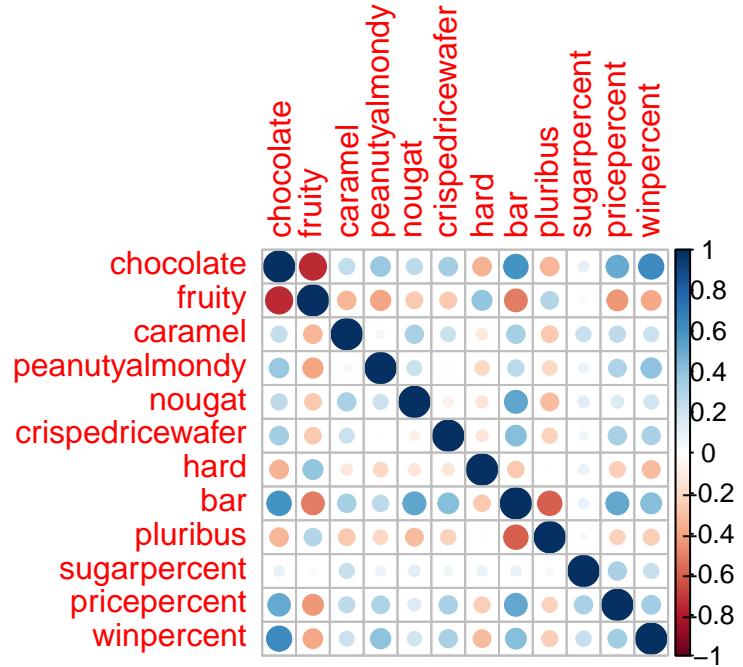
Exploring the correlation structure

Now that we've explored the dataset a little, we'll see how the variables interact with one another. We'll use `correlation` and view the results with the `corrplot` package to plot a correlation matrix.

```
library(corrplot)
```

```
corrplot 0.95 loaded
```

```
cij <- cor(candy)
corrplot(cij)
```



cij

	chocolate	fruity	caramel	peanutyalmondy	nougat
chocolate	1.0000000	-0.74172106	0.24987535	0.37782357	0.25489183
fruity	-0.7417211	1.00000000	-0.33548538	-0.39928014	-0.26936712
caramel	0.2498753	-0.33548538	1.00000000	0.05935614	0.32849280
peanutyalmondy	0.3778236	-0.39928014	0.05935614	1.00000000	0.21311310
nougat	0.2548918	-0.26936712	0.32849280	0.21311310	1.00000000
crispedricewafer	0.3412098	-0.26936712	0.21311310	-0.01764631	-0.08974359
hard	-0.3441769	0.39067750	-0.12235513	-0.20555661	-0.13867505
bar	0.5974211	-0.51506558	0.33396002	0.26041960	0.52297636
pluribus	-0.3396752	0.29972522	-0.26958501	-0.20610932	-0.31033884
sugarpercent	0.1041691	-0.03439296	0.22193335	0.08788927	0.12308135
pricepercent	0.5046754	-0.43096853	0.25432709	0.30915323	0.15319643
winpercent	0.6365167	-0.38093814	0.21341630	0.40619220	0.19937530
	crispedricewafer	hard	bar	pluribus	
chocolate	0.34120978	-0.34417691	0.59742114	-0.33967519	
fruity	-0.26936712	0.39067750	-0.51506558	0.29972522	
caramel	0.21311310	-0.12235513	0.33396002	-0.26958501	
peanutyalmondy	-0.01764631	-0.20555661	0.26041960	-0.20610932	
nougat	-0.08974359	-0.13867505	0.52297636	-0.31033884	
crispedricewafer	1.00000000	-0.13867505	0.42375093	-0.22469338	

hard	-0.13867505	1.00000000	-0.26516504	0.01453172
bar	0.42375093	-0.26516504	1.00000000	-0.59340892
pluribus	-0.22469338	0.01453172	-0.59340892	1.00000000
sugarpercent	0.06994969	0.09180975	0.09998516	0.04552282
pricepercent	0.32826539	-0.24436534	0.51840654	-0.22079363
winpercent	0.32467965	-0.31038158	0.42992933	-0.24744787
	sugarpercent	pricepercent	winpercent	
chocolate	0.10416906	0.5046754	0.6365167	
fruity	-0.03439296	-0.4309685	-0.3809381	
caramel	0.22193335	0.2543271	0.2134163	
peanutyalmondy	0.08788927	0.3091532	0.4061922	
nougat	0.12308135	0.1531964	0.1993753	
crispedricewafer	0.06994969	0.3282654	0.3246797	
hard	0.09180975	-0.2443653	-0.3103816	
bar	0.09998516	0.5184065	0.4299293	
pluribus	0.04552282	-0.2207936	-0.2474479	
sugarpercent	1.00000000	0.3297064	0.2291507	
pricepercent	0.32970639	1.0000000	0.3453254	
winpercent	0.22915066	0.3453254	1.0000000	

Q22. Examining this plot what two variables are anti-correlated (i.e. have minus values)?

Chocolaty and Fruity, bar and pluribus. etc.

Q23. Similarly, what two variables are most positively correlated?

Chocolaty and bar, fruity and hard, etc.

Principle Componenet Analysis

The function to use is called `prcomp()` with an optional `scale=T/F` argument.

```
pca <- prcomp(candy, scale=TRUE)
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	2.0788	1.1378	1.1092	1.07533	0.9518	0.81923	0.81530
Proportion of Variance	0.3601	0.1079	0.1025	0.09636	0.0755	0.05593	0.05539
Cumulative Proportion	0.3601	0.4680	0.5705	0.66688	0.7424	0.79830	0.85369
	PC8	PC9	PC10	PC11	PC12		

```

Standard deviation      0.74530 0.67824 0.62349 0.43974 0.39760
Proportion of Variance 0.04629 0.03833 0.03239 0.01611 0.01317
Cumulative Proportion  0.89998 0.93832 0.97071 0.98683 1.00000

```

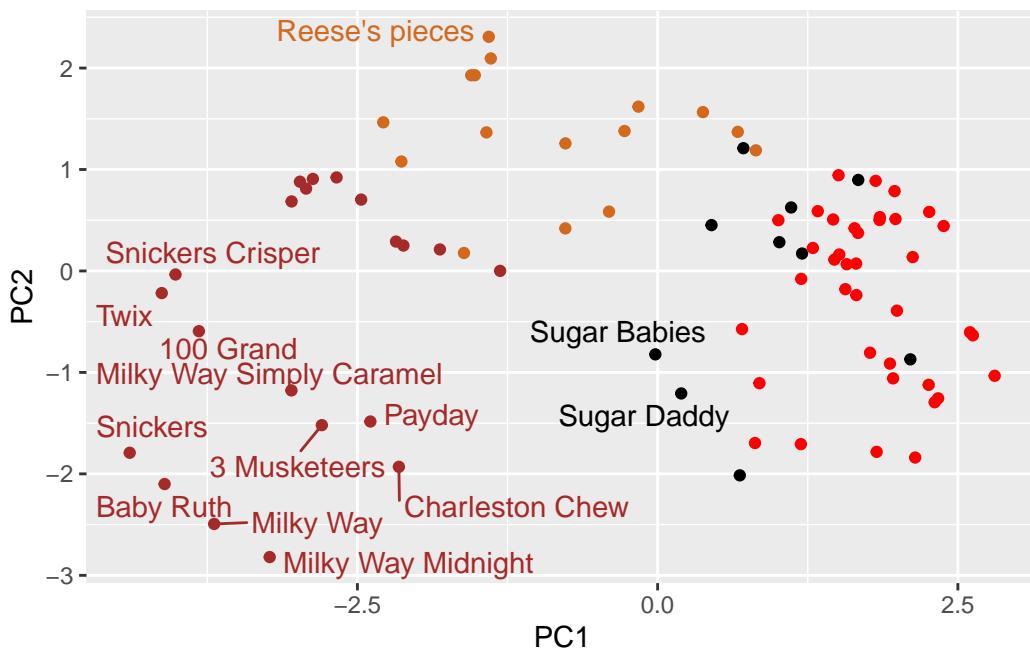
Our main PCA result figure

```

ggplot(pca$x, aes(PC1, PC2)) +
  geom_point(col=my_cols) +
  geom_text_repel(aes(label = rownames(pca$x)), col = my_cols, max.overlaps = 5)

```

Warning: ggrepel: 71 unlabeled data points (too many overlaps). Consider increasing max.overlaps



Interactive plots that can be zoomed on and “brushed” over can be made with the **plotly** package, It’s output is interactive and will not render to pdf :-)

```

#install.packages("plotly")
library(plotly)

```

Attaching package: 'plotly'

```
The following object is masked from 'package:ggplot2':
```

```
last_plot
```

```
The following object is masked from 'package:stats':
```

```
filter
```

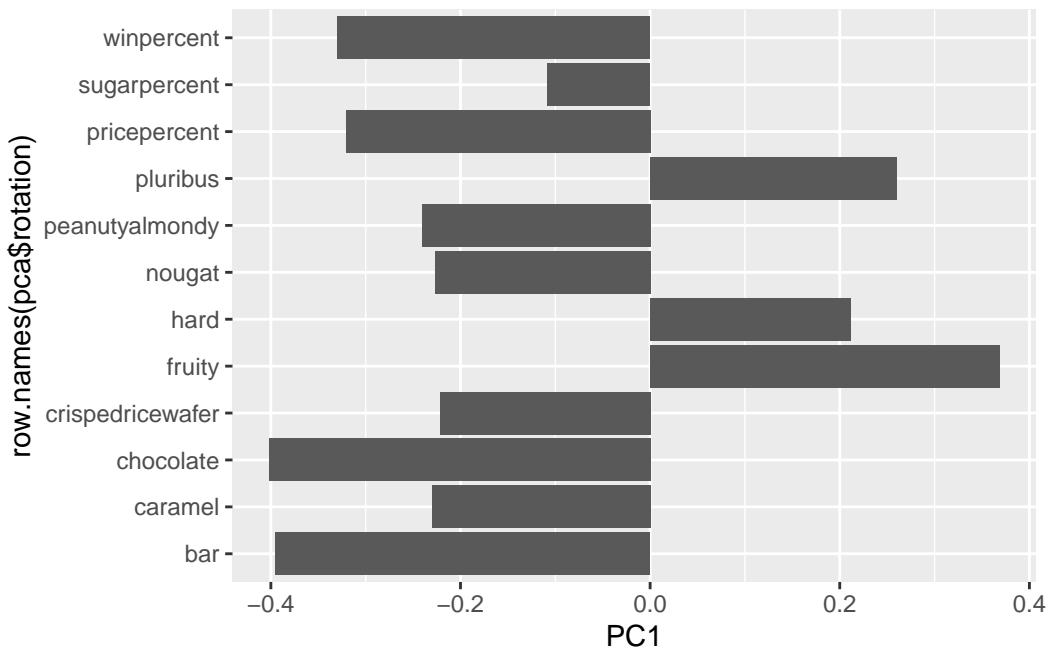
```
The following object is masked from 'package:graphics':
```

```
layout
```

```
#plotly(p)
```

We should also examine the variable “loadings” or contributions of the original variables to the new PCs

```
ggplot(pca$rotation)+  
  aes(PC1, row.names(pca$rotation))+  
  geom_col()
```



Q24. What original variables are picked up strongly by PC1 in the positive direction? Do these make sense to you?

fruity, pluribus, and hard load strongly positive on PC1.

PC1 is separating non-chocolate/fruit-style, hard, multipack candies on the positive side from chocolate/bar/price/winpercent features on the negative side—matching the known opposition between chocolate and fruity and the similarity between chocolate and bar.